AR-understanding: training and validation loss

Qwen3-8B + LoRA | 3 epochs | 36,362 training / 4,501 validation examples

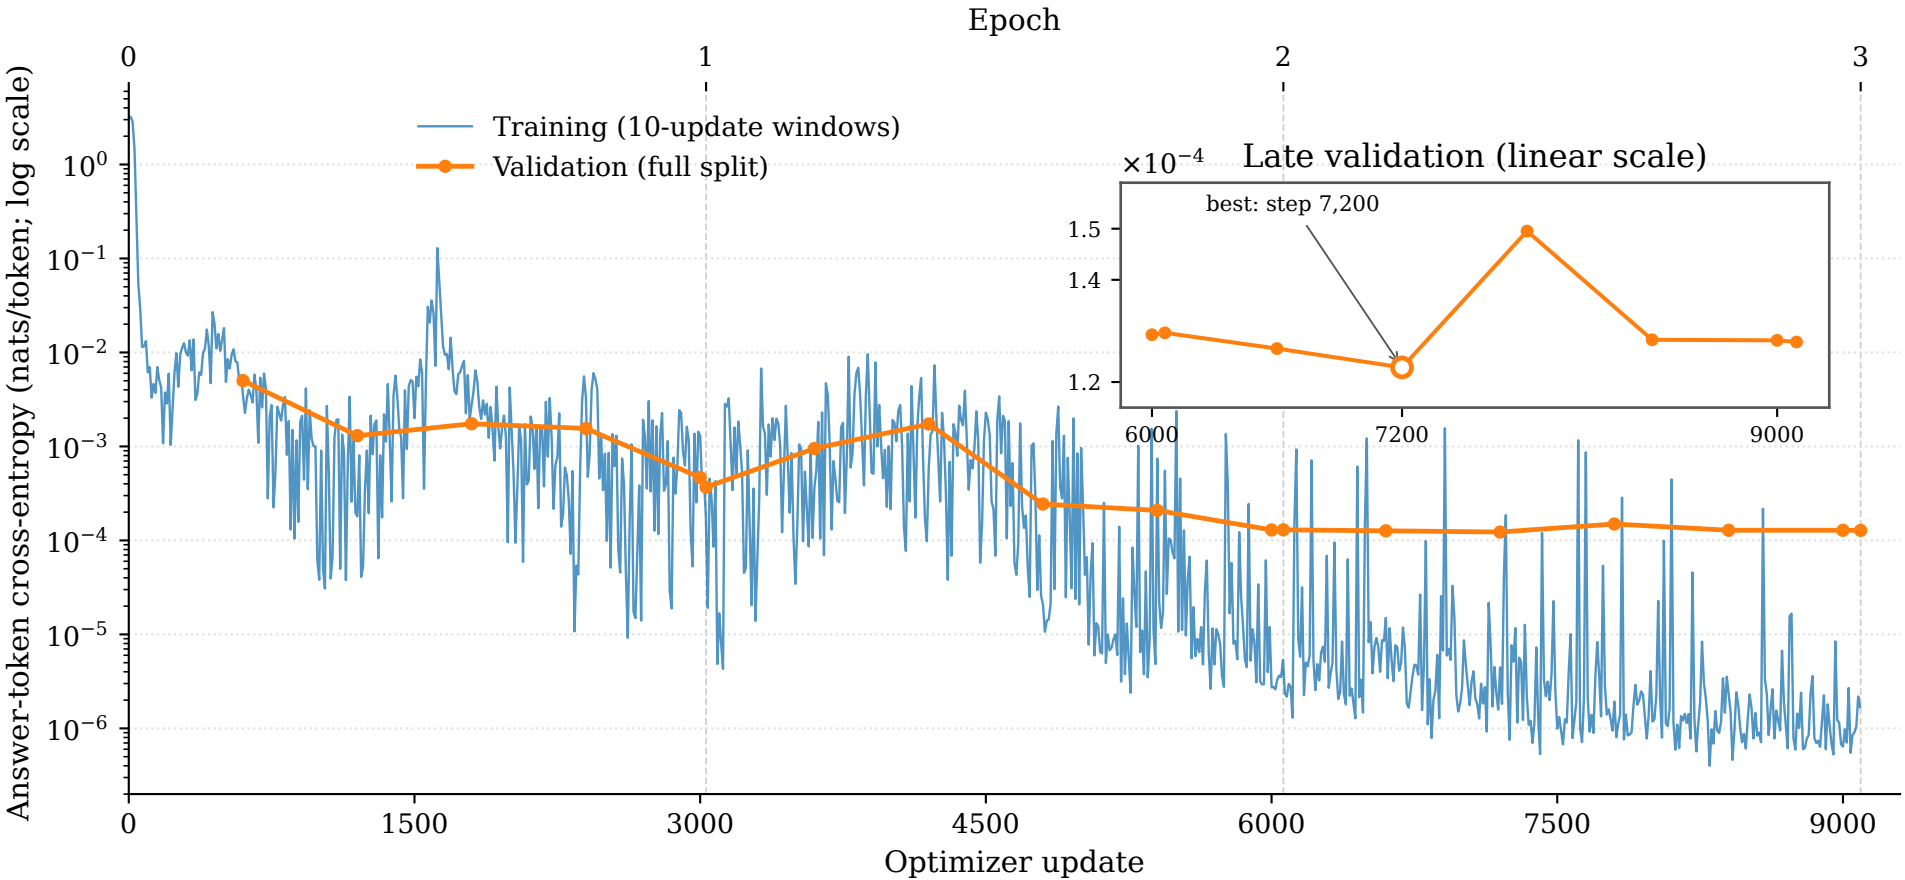

Raw recorded values; no smoothing. Prompt and padding tokens are excluded.

Last training: 1.684e-06 (step 9,090) Best validation: 1.228e-04 Final validation: 1.278e-04